

# KLEE Symbolic Execution Engine in 2019

## Authors:

Cristian Cadar, Martin Nowack

## Objective:

The objective of this paper is to provide a detailed examination of the KLEE symbolic execution engine, focusing on its development, applications, and impact since its inception at Stanford University. The paper highlights KLEE's role in advancing the field of dynamic symbolic execution, its utility in generating high-coverage tests, and its widespread adoption in both academic research and industry. Additionally, the paper discusses the evolution of KLEE's architecture, its community contributions, and its performance in competitive benchmarks, emphasizing the tool's adaptability and extensibility in various domains such as automated debugging, exploit generation, and software verification.

## Concept / Methods / Techniques:

KLEE operates on the principle of dynamic symbolic execution (DSE), which involves systematically exploring multiple execution paths in a program by treating inputs as symbolic values rather than concrete ones. This technique allows KLEE to automatically reason about the feasibility of different paths using constraint solvers, making it possible to uncover bugs and generate comprehensive test cases. The engine works at the level of LLVM bitcode, enabling it to analyze a broad spectrum of programs across different languages. KLEE's modular and extensible architecture supports a variety of search heuristics, constraint solvers, and optimization techniques, which can be easily customized or extended according to specific testing needs. This flexibility has made KLEE a versatile tool in the field of software testing and analysis, with applications extending beyond traditional test generation to include tasks like program repair, fault reproduction, and patch analysis.

## Data:

The paper provides evidence of KLEE's significant impact and widespread use in both academic and industrial settings. Since its migration to GitHub in 2013, KLEE has attracted contributions from over 60 developers, reflecting the collaborative nature of its development. The mailing list associated with KLEE has garnered over 350 subscribers, indicating a strong community of users and developers actively engaging with the tool. KLEE's original paper has been cited over 2000 times, showcasing its influence in the field of symbolic execution. Furthermore, KLEE was featured prominently in the First International KLEE Workshop on Symbolic Execution in 2018, which saw participation from over 80 attendees from academia, industry, and government. This event, along with the ongoing contributions from the KLEE community, underscores the tool's relevance and the continued interest in its development and application.

### **Abstract:**

KLEE is a dynamic symbolic execution engine that originated from Stanford University and has since been primarily developed by the Software Reliability Group at Imperial College London. The tool has established itself as a key player in the field of software testing and analysis, particularly in test generation and bug finding. The paper traces the history of KLEE, highlighting its foundational concepts, the development of its architecture, and the contributions from a wide range of developers and researchers. KLEE's performance in the 2019 Test-Comp competition is also discussed, where it demonstrated its strengths in systematic exploration and constraint solving, despite challenges such as path explosion. The paper also emphasizes KLEE's adaptability, with its modular design allowing users to tailor the tool to their specific testing requirements, thereby making it a valuable resource for a variety of applications.

### **Conclusion:**

KLEE continues to be a significant tool in the domain of symbolic execution and software testing. Its ongoing development, driven by both academic and industrial contributors, has ensured its relevance and effectiveness in identifying software bugs and generating test cases. While challenges such as path explosion and the complexity of constraint solving persist, the paper suggests that future efforts should focus on addressing these issues and enhancing KLEE's applicability to larger, real-world software systems. The successful participation of KLEE in competitive benchmarks like Test-Comp 2019 highlights its robustness and the potential for further refinement. The paper concludes by acknowledging the strong community support that has been crucial to KLEE's development and suggesting that continued collaboration will be key to overcoming the challenges that remain in the field of dynamic symbolic execution.